

PROPRIETARY FRAMEWORK

GOVERN FIRST

Building an AI Centre of Excellence that survives past the first pilot

Every AI CoE fails the same way — for opposite-looking reasons

The Bottleneck

Governance-first, nothing ships

- Every use case routes through committee
- Risk sign-off takes months, not weeks
- Business units quietly build their own tools instead

The Pilot Factory

Delivery-first, nobody adopts

- Flashy demos win the room, then stall
- No governance trail when Risk finally notices
- First real failure gets the whole programme pulled

THE CORE INSIGHT

Both get built by capable people. Both fail for the same reason:

The sequence is backwards.

Governance built after delivery reads as damage control. Built before it, it's what lets a CoE move fast without being the reason it gets shut down.

Four layers, applied in order — skipping ahead is the failure mode

01	Operating Model	Decide the shape of the CoE, centralised, enablement, or hybrid, before you build anything.
02	Governance	Build the approval scaffolding before the first use case, not after it ships.
03	Use Case Portfolio	Sequence opportunities by exposure and value, not by which one looks best in a demo.
04	Capability & Monitoring	Train the organisation and monitor for drift and fairness in parallel with delivery, not after.

Operating Model

Decide how the CoE relates to the rest of the business before you build anything

01 Regulatory exposure How much of the likely use case surface touches regulated decisions or protected data.

02 Existing capability How much AI-building skill already exists inside business units, outside the CoE.


03 Spend model Whether budget and headcount sit centrally or are distributed across business unit P&Ls.

Governance

The scaffolding gets built before the first use case, not after it ships

01	Regulatory exposure	Does this use case touch a regulated decision or a customer-facing output with harm potential?
02	Data sensitivity	What data does the model see, and where does it sit once it's processed?
03	Reversibility	If the output is wrong, how easily and cheaply is it corrected, before or after a customer sees it?

Sequence use cases by exposure, not just value




The highest-value use case is rarely the right one to ship first. The right one is the lowest-exposure use case that still proves the model works.

Capability & Monitoring

Training and monitoring run alongside delivery, not after it


01	Capability building	Role-based training and a community of practice, started before the first use case is even visible.
02	Value monitoring	Benefit tracked against baseline every month, owned by a named individual, not a dashboard nobody reads.
03	Fairness & audit	Output checked across segments for drift, with a human override path and a reconstructable audit trail.

Four stages, twelve to eighteen months



Not one-size-fits-all, and that's deliberate

The four layers hold regardless of industry. What changes is how much formality each layer needs to carry.



Illustrative, not derived from a named study — the scoring worksheet (Layer 1) produces the real number per organisation.

Regulated & large

Named sign-off from Risk and Compliance, formal intake gate, hybrid operating model centralised to start.

Smaller & lower exposure

One person exercising judgement against the same three questions, gate still exists, just not a committee.

What the first eighteen months look like

Days 0–90

Operating model decided. Governance gate built. First 1–2 low-exposure use cases resourced.

Days 90–180

First use case live with a measured benefit. Monitoring dashboard stood up. Training launched.

Months 12–18

Portfolio of 8–12 use cases. At least one business unit self-serving under CoE oversight.

WHERE THIS GETS SPECIFIC

This is the map. The build is the territory.

The scoring that decides which operating model fits. The exact questions a use case has to answer before it gets resourced. The grid that sequences a real portfolio against real regulatory exposure. That's where this stops being a framework and starts being organisation-specific work.



Abhinav

AI strategy and governance leader, working across CPG, Retail, Luxury and Financial Services on AI roadmaps, adoption, and responsible-use frameworks. Author of Pixelated Souls and developer of a clinical psychological framework used in supervised practice.

- AI strategy and CoE build across regulated and unregulated sectors

- Governance frameworks that hold up to Risk, Compliance and Information Security scrutiny

- Bringing a psychologist's read on adoption and behaviour change to a technical build

BUILDING AN AI CoE IN YOUR ORGANISATION?

Let's talk sequence.

I help organisations design the operating model, governance gate and use case portfolio that let AI adoption move fast without losing the trust of Risk, Compliance and the board.

[Email placeholder]

[LinkedIn placeholder]