

PROPRIETARY FRAMEWORK

SOUNDINGS

Knowing what's running, and whether it's watched properly, before you build anything

Abhinav — AI Strategy & Governance

Private & confidential · © All rights reserved

Nobody can answer a simple question

Ask five people in the same company whether AI is being used correctly and you get five different answers, because nobody agreed what correct means.

1

No inventory

Most companies can name a handful of sanctioned tools. Real usage runs into the hundreds, most of it never approved by anyone.

2

No standard

There is no single right way to use AI. A chatbot drafting an email and a model scoring a loan application need completely different amounts of control.

3

No timing

Most companies find out what's actually running only after something breaks, not before.

The failure isn't the wrong AI. It's the wrong amount of oversight for what it actually does.

High-risk work running with no one watching it. Low-risk work buried under process so heavy that people quietly build their own tools to get round it. Both are the same problem, pointing in opposite directions.

Score exposure and oversight separately

RISK

Regulatory exposure

Does this touch a regulated decision, or a customer-facing output with harm potential?

Data sensitivity

What does the model see, and where does it sit once it's processed?

Reversibility

If the output is wrong, how easily and cheaply is it corrected?

OVERSIGHT

Governed

Sanctioned, reviewed, actively monitored.

Informal

Known about. Nobody is actually checking.

Invisible

Nobody in the organisation has agreed it exists.

Every answer comes from a fixed checklist, not a rating scale. Same inputs, same score, regardless of who is running the audit.

Plot exposure against oversight and the danger sorts itself



A structured read, not a network scan

1

Amnesty first

Teams hear upfront that disclosing an unsanctioned tool isn't a disciplinary conversation. Without that, people underreport, and the map is wrong before it starts.

2

Structured interviews and document review

Every AI use surfaced is scored against the same risk and oversight criteria, use case by use case, not tool by tool.

3

Delivered in days

No log access, no SSO integration, no procurement cycle. A sounding is fast by design, that's what makes it a sounding and not a survey.

Three things land on the table

The mismatch map

Every use case plotted. The exposure zone named and ranked, not buried in an appendix.

A prioritised action list

Sequenced by what's actually dangerous, not by what's loudest in the room.

The inputs Keel needs

The operating model and governance gate decisions, made on evidence instead of guesswork.

A sounding is taken before the ship commits to open water

Keel is the structural spine, decided once you know what you're building. Soundings is the depth check that tells you whether the water's safe to move into. One measures. One builds. Skipping the measurement is how a CoE ends up governing use cases it never actually understood.

SOUNDINGS

Measure. Days, not months.



KEEL

Build. Twelve to eighteen months.

Same instrument, two ways in

INTERNAL

Run inside a company that employs you.

The first visible move in the first weeks of a new role or a new mandate. Becomes the evidence base for why a full CoE build is actually needed, and what shape it should take.

EXTERNAL

Sold as a fixed-fee, short engagement.

Stands alone for a company not ready for a full CoE. Converts straight into a Keel engagement for one that is. One diagnostic, two commercial doors.

ABOUT THE AUTHOR

Abhinav

AI strategy and governance leader, working across CPG, Retail, Luxury and Financial Services on AI roadmaps, adoption, and responsible-use frameworks. Author of Pixelated Souls and developer of a clinical psychological framework used in supervised practice.

- AI strategy and CoE build across regulated and unregulated sectors
- Governance and audit frameworks that hold up to Risk, Compliance and Information Security scrutiny
- A psychologist's read on why people hide AI use, not just where it's hiding

WANT TO KNOW WHAT'S ACTUALLY RUNNING?

Let's take a sounding.

A short, fixed-fee read on every AI use case running in your organisation, scored for exposure and oversight, delivered in days.

[Email placeholder]

[LinkedIn placeholder]